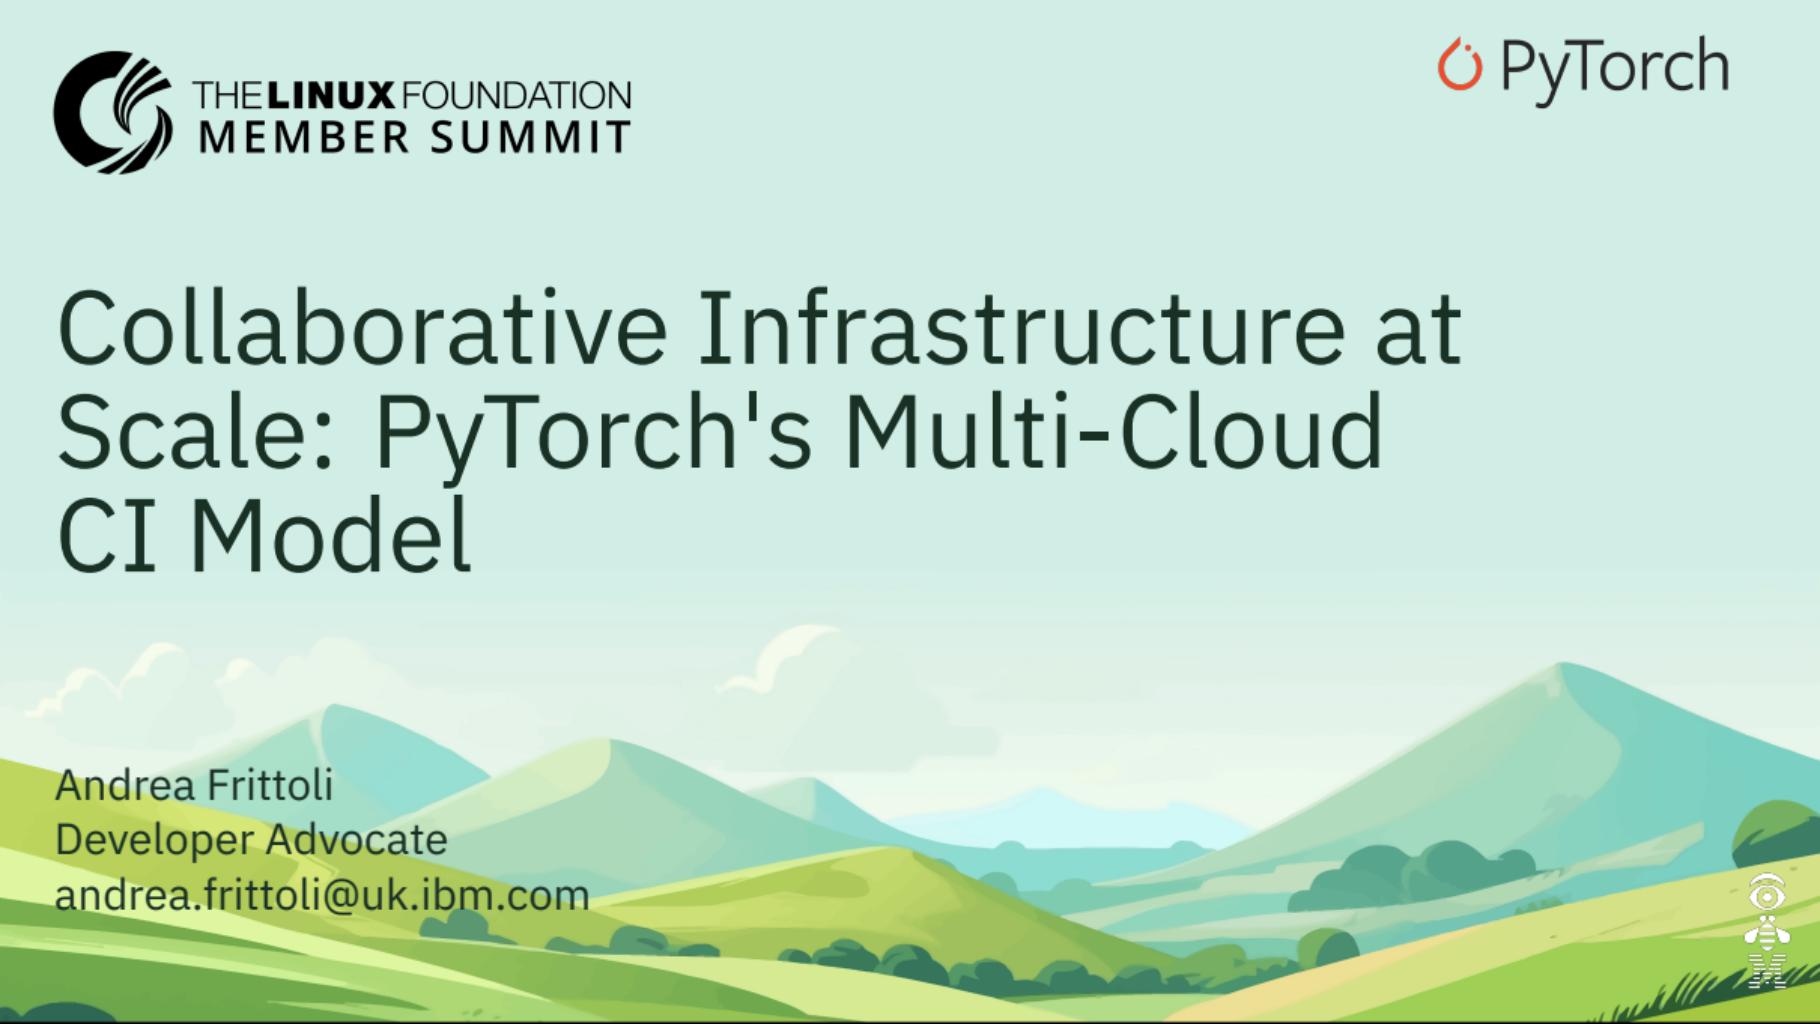


Collaborative Infrastructure at Scale: PyTorch's Multi-Cloud CI Model



Andrea Frittoli
Developer Advocate
andrea.frittoli@uk.ibm.com



Andrea Frittoli

⌚ afrittoli | 💬 andreafrittoli | 🗣 @blackchip76

- › Open Source Advocate @ IBM
- › Lives in Wales, enjoys the wind
- › Multi-Cloud CI Working Group Lead
- › Member of PyTorch Infra Working Group
- › Tekton, CDEvents maintainer





Contents

The Infrastructure Challenge
PyTorch Foundation
Governance Model
Technical Architecture
Results & Lessons Learnt
Q&A



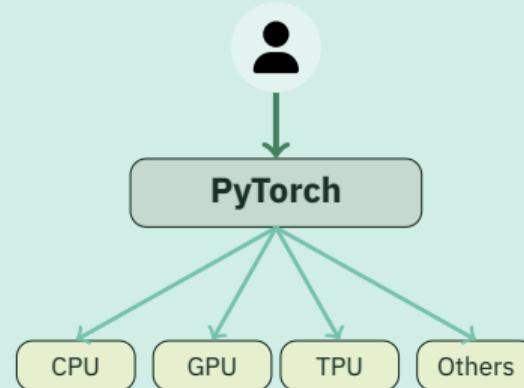
The Infrastructure Challenge



What is PyTorch?

A Leading Open Source ML Framework:

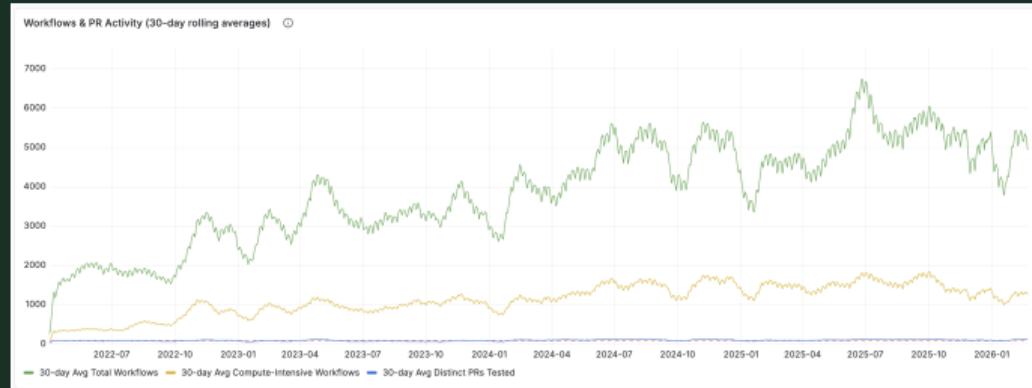
- › Researchers, data scientists, ML engineers
- › Research and production ML workloads
- › (Distributed) Training, LLM, RL, Inference
- › Python API, C++ Core (libtorch)
- › Eager and Graph Mode (Inductor)
- › Rapid growth and adoption across industry



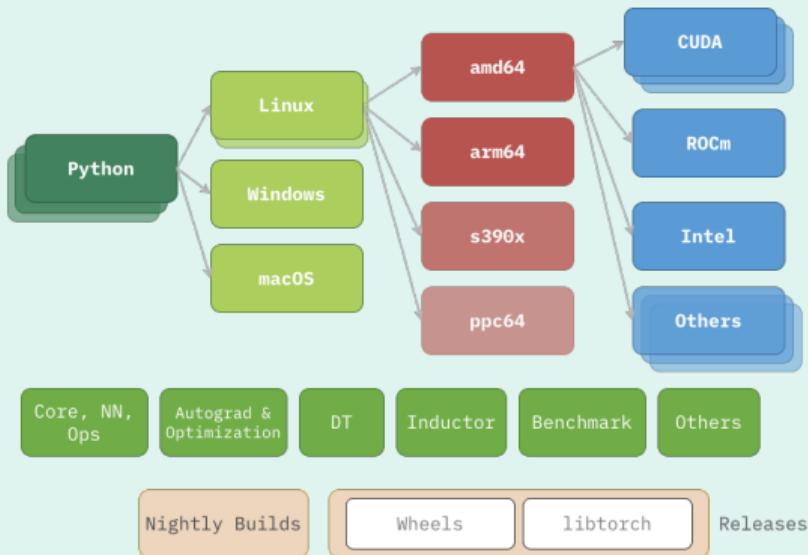
CI/CD Infrastructure Today

- › ~1M\$ monthly infrastructure costs
- › ~1M hours of compute per month
- › ~100 distinct PRs/day
- › ~1.3k Compute-intensive Workflows/day
- › Growing year over year

pytorch/pytorch repo only, numbers are estimated

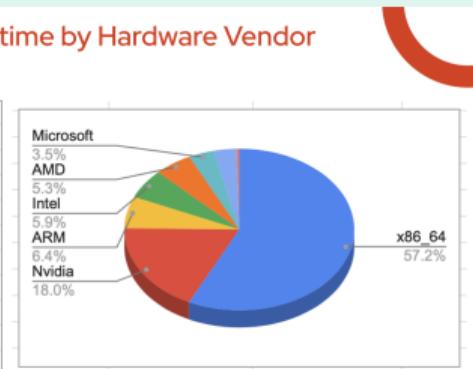


An Expanding Test Matrix



Combined Hardware Runtime by Hardware Vendor
September 2025

Combined Runtime By Hardware Vendor	
Vendor	Total
x86_64	834,085 hours
Nvidia	262,878 hours
ARM	93,000 hours
Intel	86,090 hours
AMD	77,517 hours
Microsoft	51,515 hours
Apple	46,812 hours
IBM	6,245 hours
N/A	0 hours
1,458,142 hours	



* Data from PT Foundation AWS Account + PyTorch HUD. This does not include self-hosted runner costs by member provided runners.



The Challenges

Community Challenges:

- › Maintain high-quality end-user experience
- › Preserve contributor workflow
- › Provide clear path for vendor engagement

Scale Challenges:

- › **Engineering:** Managing a Diverse Fleet, Access to Platforms
- › **Financial:** Reconcile Community Wishes with Budget Constraints

Security Challenges:

- › Trust and integrity in a diverse environment
- › Central Administration
- › Trusted Builds

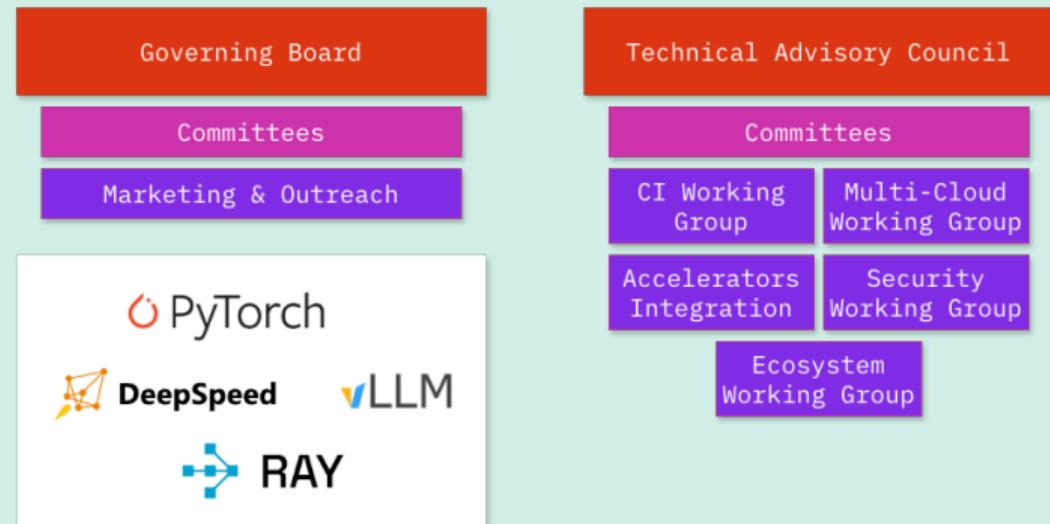


PyTorch Foundation & Working Groups



The PyTorch Foundation

- › Part of the Linux Foundation
- › Neutral home for PyTorch project
- › Hosts multiple projects
- › Thought Leadership (GB)
- › Technical Leadership (TAC)
- › Marketing & Outreach



Working Groups

Multi-Cloud CI

Cloud Agnostic Infrastructure:

- CI/CD jobs and Supporting Infra
- Metrics and Monitoring
- Fleet Management

Vendor-managed Runner Pools

CI Infra

Responsible for the CI/CD infra:

- Develop and maintain tools
- Operate the AWS/GitHub fleet
- Execute Releases

Accelerator Integration

Software Integration:

- Developer Guide
- Framework Improvement

Reference Test Framework

CI Event Relay Infrastructure

Security

- PyTorch Security Triage
- Tools and reports
- CI/CD Security



Governance Model



Photo by Khampha Phimmachak, CC0



Personas: A Balancing Act

End Users

- Stable software
- Available on my dev platform
- Available on my prod platform
- Following the industry at speed

Platform Vendors

- PyTorch for my platform
- Demonstrate platform compatibility
- Low bar to contribution of infrastructure
- Clear guidance and processes

Project Maintainers

- Project Success
- Stay relevant for end-users
- High quality secure builds
- Manageable CI/CD System (scale)
- Smooth CI/CD experience for contributors

PyTorch Foundation

- Ensure Vendor Neutrality
- Vendor participation encouraged
- Fair representation
- Financial sustainability



Integration Models

Three Ways to Contribute Infrastructure:

- › **Public Cloud Credits:** Financial contribution for shared infrastructure
- › **Vendor-Managed Runner Pools:** Tightly integrated with GitHub Actions
- › **Vendor-Managed Test Infrastructure:** Loosely integrated with GitHub

	Community	Scale	Security
Advantages	Flexible models accommodate different vendor capabilities	Outsourcing engineering with vendor-managed setups	More loosely coupled systems cannot compromise core CI system
Challenges	Hide the underlying infra complexity to contributors	Consistent tech stack required for public cloud credits. Consistent metrics and monitoring across the board.	Ensure security best practices on vendor managed infrastructure



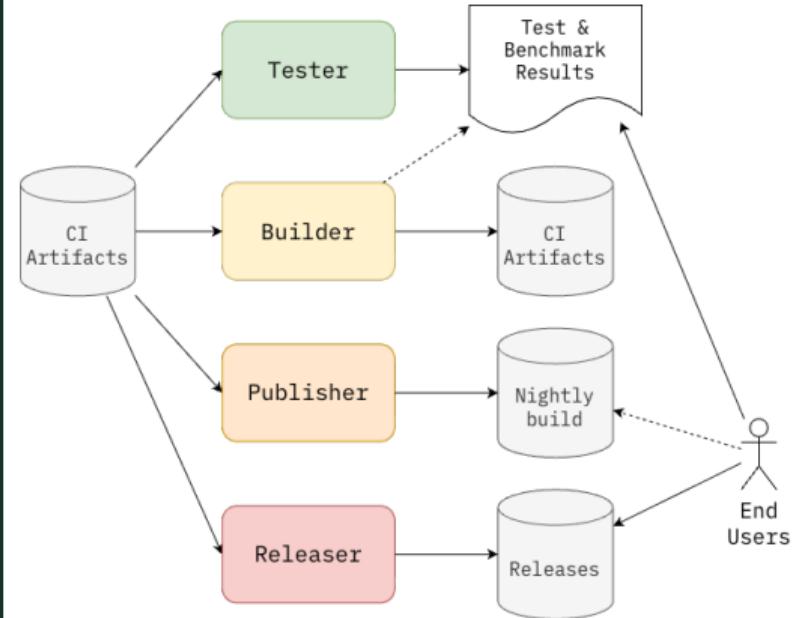
Roles & Requirements

Least privilege access to cloud resources:

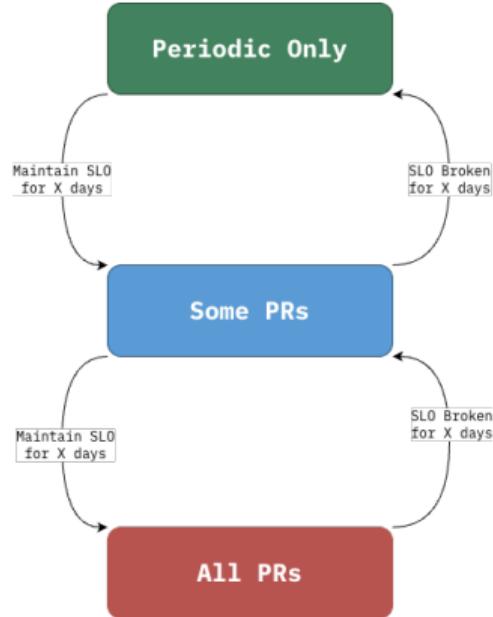
- › Jobs specify nodes access level required
- › Runners are assigned a role
- › Roles define the level of access to resources

Requirements for Runners:

- › Configuration
- › Provisioning
- › Monitoring
- › Security



Triggers & SLOs



Trigger levels for CI/CD jobs:

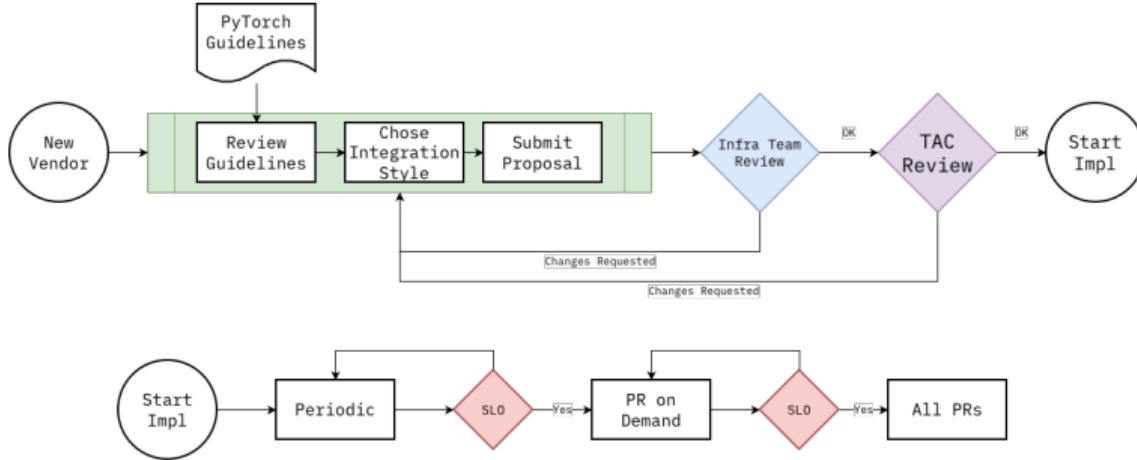
- On-demand only
- Periodic nightly
- Periodic multi-times per day
- Specific PRs only
- All PRs

Service Level Objectives:

- Availability %
- Reliability (Num of infra failures)
- Average Job Queue Time (p95/3 months)
- Engineering Commitment (on-call, slack, meetings)



Contribution Process



- › Working Group provides guidance and recommendations
- › Infra team verify checks for compliance
- › TAC Community for final approval
- › Transparent process with clear requirements
- › Monitoring for SLOs



Central Administration & Visibility

Maintaining Manageability at Scale:

- › Common software stack across clouds
Reusable by vendors too
- › Self-service Tools with Central Admin Override
- › Vendor playbooks, reusable IaaC

Visibility & Control:

- › Public Dashboards: Transparency for community and vendors
- › Cost Tracking: Attribution per vendor and workload type
- › Performance Metrics: Build times, success rates, SLA compliance

A screenshot of a GitHub Actions dashboard for the 'pytorch/pytorch' repository. The main view shows a grid of small cards, each representing a different pull request (PR) and its current status. The cards include information like the PR number, author, and a color-coded icon indicating the status (e.g., green for success, red for failure). The grid is organized by time, with the most recent PRs at the top.

A screenshot of the GitHub Actions Runner Service dashboard. It features a summary section with four colored boxes: blue for 'Total Runners' (4), green for 'Active' (0), grey for 'Offline' (4), and yellow for 'Pending' (0). Below this, there's a 'Recent Activity' section listing three recent events: 'batch_delete_runner' by admin@githubs.org, 'batch_delete_runner' by admin@githubs.org, and 'label_policy_update' by alice@githubs.org. Each event is timestamped and includes a link to view more details.

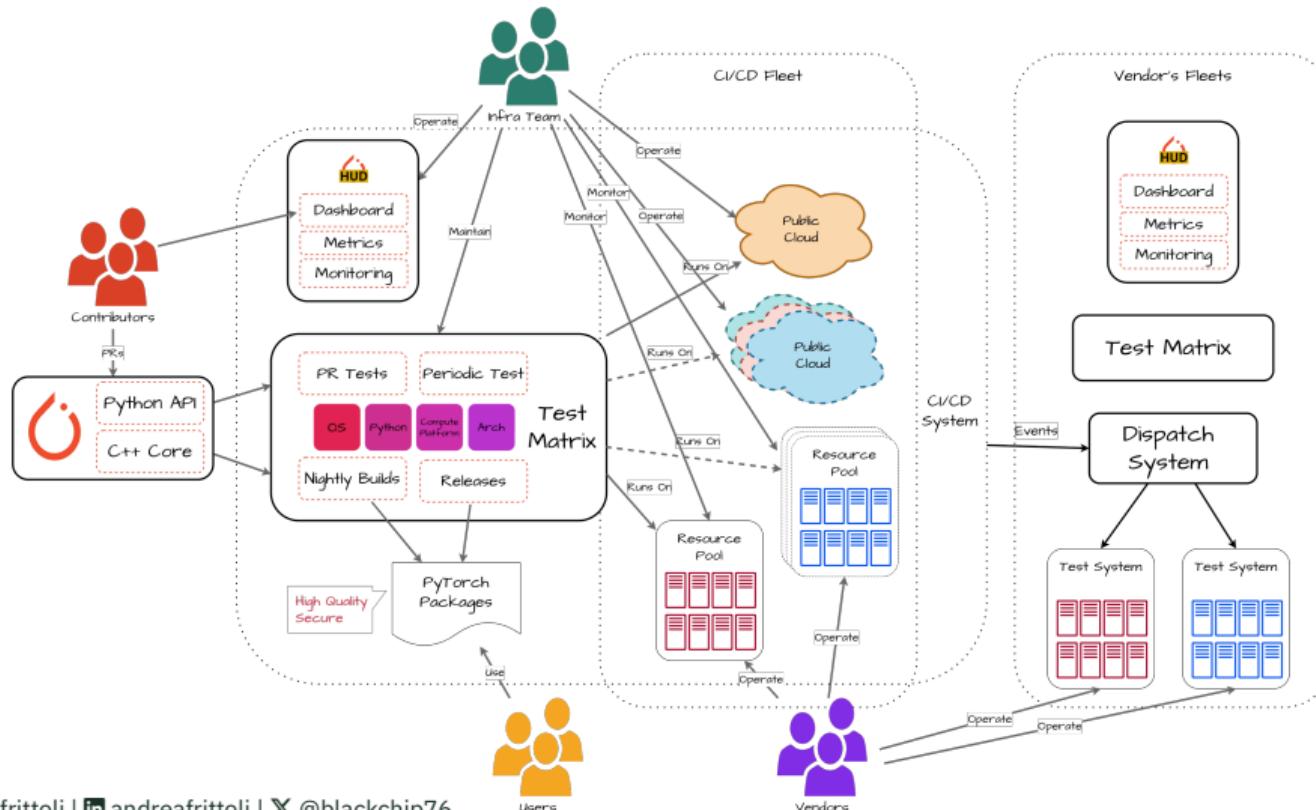
A screenshot of the GitHub Actions Runner Service dashboard under the 'Security Events' tab. It lists several recent events, each with details like timestamp, event type, severity, actor, and target. The events shown are: 'batch_delete_runner' (high severity, admin@githubs.org, target -) on Jan 26, 2026 at 5:09 PM; 'batch_delete_runner' (high severity, admin@githubs.org, target -) on Jan 26, 2026 at 5:08 PM; 'label_policy_update' (medium severity, alice@githubs.org, target alice-number-dash12) on Jan 26, 2026 at 5:40 PM; and 'batch_delete_runner' (high severity, admin@githubs.org, target -) on Jan 26, 2026 at 5:32 PM.



Technical Architecture



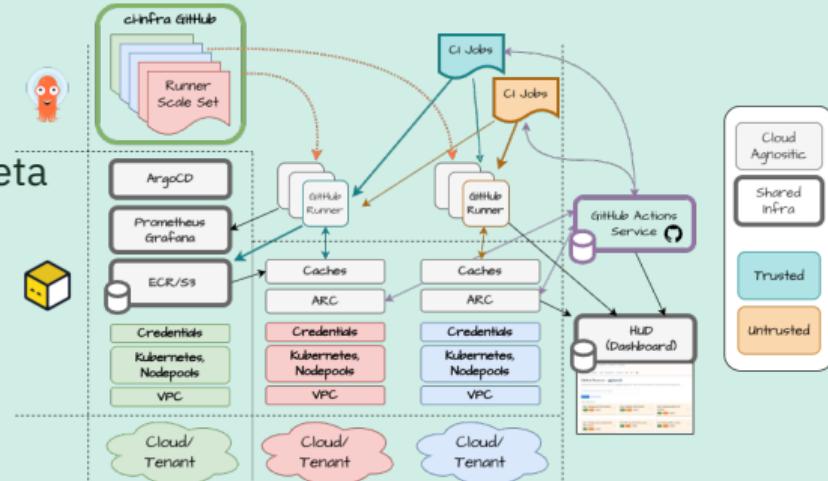
High Level Architecture



Public Cloud Infrastructure

Current AWS Setup:

- > Primary infrastructure hosted on AWS
- > Two accounts, funded by AWS Credits and Meta
- > Self-hosted GitHub Actions runners
- > Managed by PyTorch Infra team
- > Autoscaling based on workload demand



Multi-Cloud & Infra Working Group Initiatives:

- > **Portable CI Jobs**: Run in a container, remove cloud-specific assumptions
- > **Reusable Autoscaler**: Developing portable autoscale based on ARC
- > **Infrastructure as Code**: Reusable OpenTofu modules



Vendor-Managed Runner Pools

Challenges:

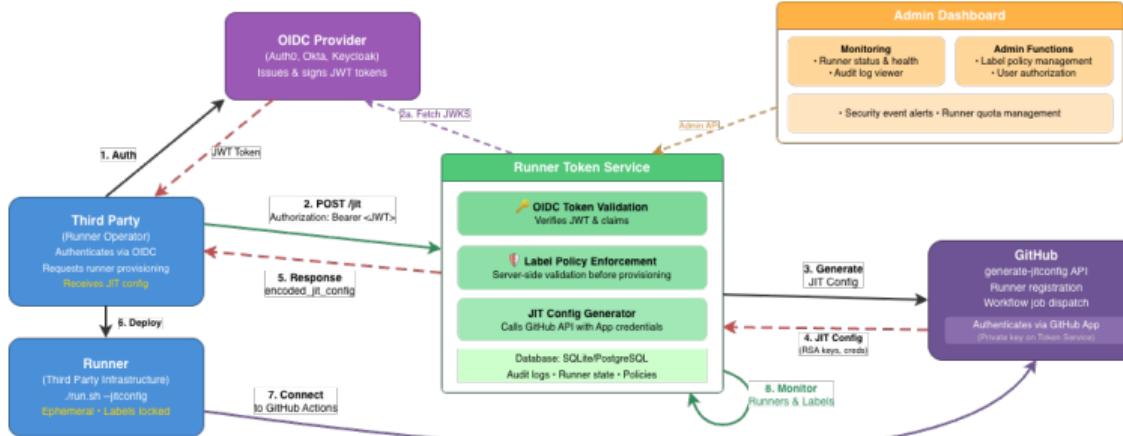
- › Long lived admin credentials
- › Runner metadata cannot be enforced
- › Lack of central administration
- › Lack of auditability

Proposed Solution:

- › GHA Runner Token Service (GHARTS)
- › Use LFID for authentication
- › No Admin Credentials to vendors
- › Ephemeral, non reusable credentials
- › Enforcement of metadata and quota
- › Centralized administration
- › Auditability



GHARTS Provisioning Flow



- › Vendor authenticates using existing credentials (LFID)
- › Vendor requests a JIT config using the JWT
- › GHARTS verifies AuthN and AuthZ for vendor
- › GHARTS requests JIT config from GitHub
- › Vendor uses JIT config to provision a runner
- › JIT lasts 1h, and can only be used Once



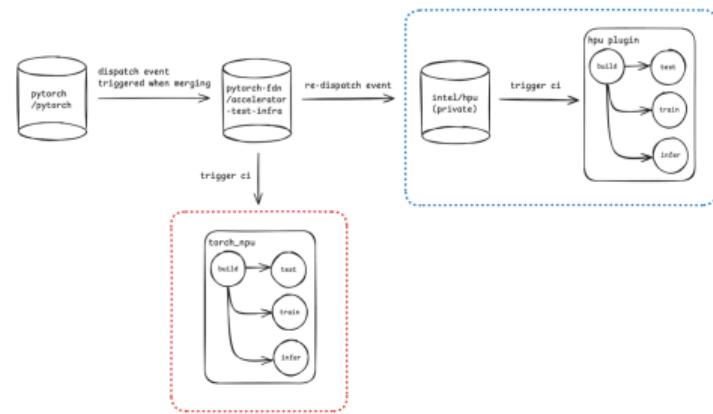
Vendor-Managed Test Systems

Accelerator Integration Working Group Initiative:

- Loosely coupled with PyTorch CI
- Events relayed to external repository
- Vendors provide test infrastructure and environments
- Out-of-tree test workflows

Use Cases:

- Specialized hardware testing (TPUs, custom accelerators)



Old Issues, New Issues

PyTorch Queue Time Analysis

* All datetime values are in UTC. Local: 2026-02-24 21:36:29. UTC Time: 2026-02-25 05:36:29

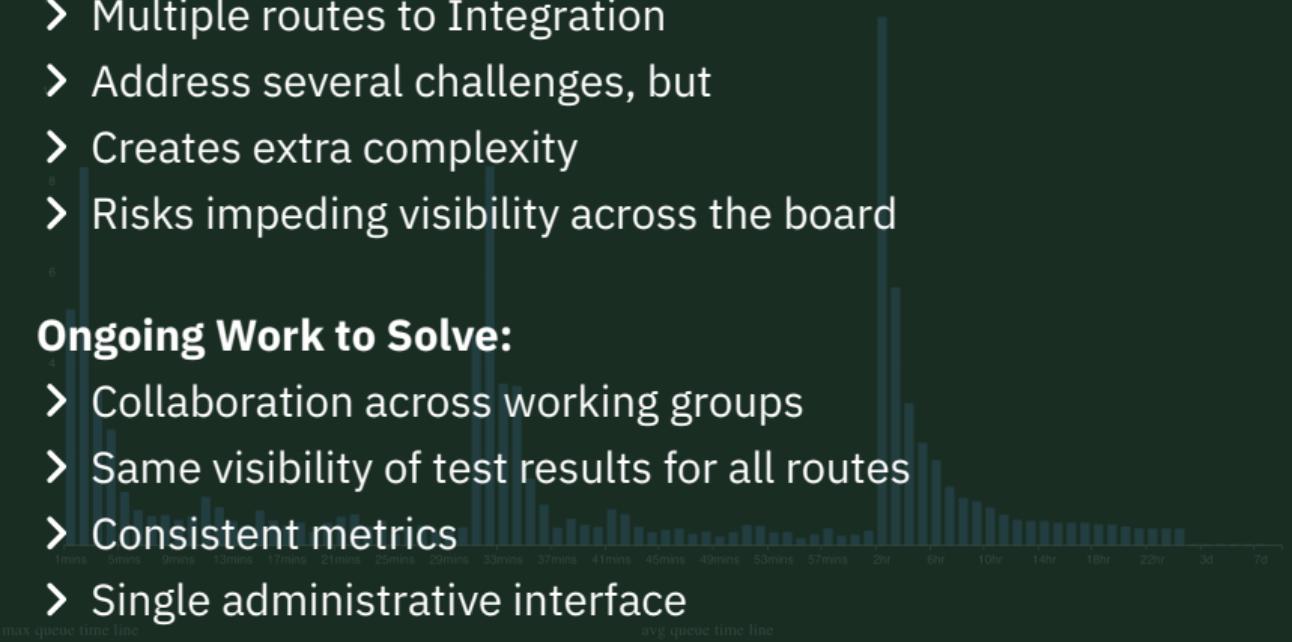
* Data is collected every 30 minutes, including all jobs in queue at that time. ⓘ

Chart Type



New Challenges:

- Multiple routes to Integration
- Address several challenges, but
- Creates extra complexity
- Risks impeding visibility across the board



Ongoing Work to Solve:

- Collaboration across working groups
- Same visibility of test results for all routes
- Consistent metrics
- Single administrative interface

max queue time line

avg queue time line

Time Range Last 3 Days Granularity Half Hour

Search Category workflow name

By default, shows data for all queued jobs. Using filter and checkbox below for specific items.

Select items:

Filter

Select All Unselect All

- pull
- inductor
- operator_microbenchmark
- inductor-A100-perf-nightly
- inductor-perf-nightly-h100
- windows-binary-litbtorch-debug
- trunk
- linux-aarch64-binary-manywheel
- periodic
- inductor-unitest
- windows-arm64-binary-wheel
- trunk-rocm-sandbox
- windows-binary-wheel

Search

Click to apply filter changes



Results & Lessons Learnt



Fleet Composition

Onboarded:

- › **Public Cloud Credits:** AWS, Meta
- › **Vendor-Managed Runner Pools:** AMD ROCm, IBM s390x, IBM ppc64, Intel XPU, NVIDIA CUDA B200, NVIDIA Windows RTX
- › **Vendor-Managed Test Systems:** Huawei Ascend NPUs, Intel Gaudi HPU
- › ...and GitHub Hosted runners too!

In progress:

- › **Vendor-Managed Runner Pools:** IBM s390x (via GHARTS)
- › **Vendor-Managed Test Systems:** RedHat RHEL



Ongoing Work

Public Clouds

- Migrate CI Jobs (with the help of AI!)
- Migrate Autoscaler to ARC based system

Runner Pools

- Deploying GHARTS
- Onboarding IBM pools
- Integrate ARC with GHARTS

Test Systems

- Integrate into hud
- Enable PR comments and votes

Governance

- Finalize onboarding process
- Finalize SLO definitions



Lessons Learned

- > **Clear onboarding:** Streamlined process and transparent decision making
- > **Document SLOs:** Define roles and SLOs, include community engagement
- > **Reusable Tech Stack:** Share proven configurations and best practices
- > **Engage Vendors in WG:** Vendors can and want to contribute to the solution
- > **Self Service First:** Distributed engineering effort
- > **Empower Maintainers:** Give infra team tools to manage at scale
- > **Security:** Consider contributed code and infrastructure as untrusted by default
- > **No one-size-fit-all:** Multiple contribution paths reduce attrition



What's Next

- › Third-party builds and releases
- › Consistent metrics, monitoring and UX across the fleet
- › Adopt standards like OpenTelemetry/CDEvents
- › Continue balancing community needs and project sustainability



To summarize

- PyTorch builds the LLMs that power the AI revolution
- Hosted by the PyTorch Foundation

CI/CD Infra challenges:

- community
- scale (financial/engineering)
- security

- The solution is a mixture of governance and technology
- Engaged with Vendor through Working Groups to address challenges

- Multiple paths to contribution, consolidate visibility and administration
- Work continues to make the CI/CD truly cloud agnostic





Questions?



References & Contact

Resources:

- › PyTorch: <https://pytorch.org>
- › PyTorch GitHub: <https://github.com/pytorch/pytorch>
- › PyTorch Foundation: <https://pytorch.org/foundation>
- › Multi-Cloud Infrastructure WG:
<https://github.com/pytorch-fdn/multicloud-ci-infra>
- › Accelerator Integration WG:
<https://github.com/pytorch-fdn/accelerator-integration-wg>



Contact:

- › Andrea Frittoli
- › [LinkedIn](#) andreafrittoli | [GitHub](#) afrittoli

