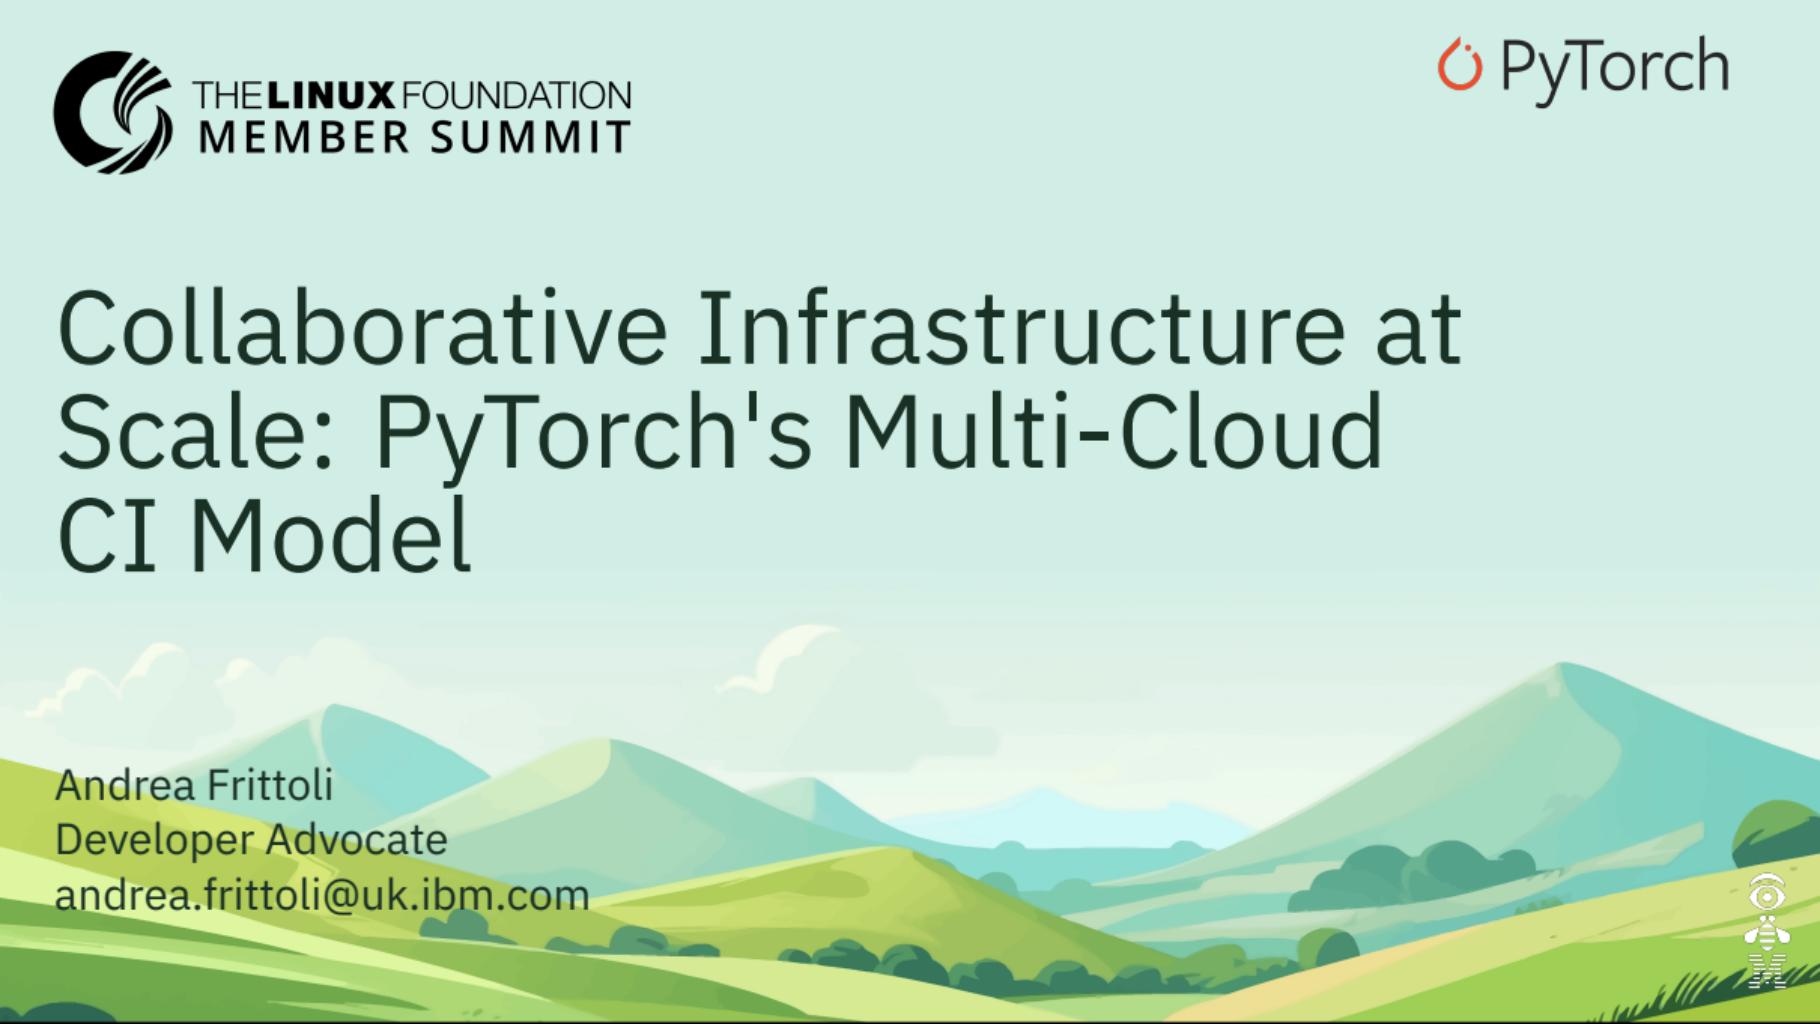


Collaborative Infrastructure at Scale: PyTorch's Multi-Cloud CI Model



Andrea Frittoli
Developer Advocate
andrea.frittoli@uk.ibm.com



Andrea Frittoli

⌚ afrittoli | 💬 andreafrittoli | 🗣 @blackchip76

- › Open Source Advocate @ IBM
- › Lives in Wales, enjoys the wind
- › Multi-Cloud CI Working Group Lead
- › Member of PyTorch Infra Working Group
- › Tekton, CDEvents maintainer





Contents

The Infrastructure Challenge

PyTorch Foundation

Governance Model

Technical Architecture

Results & Next Steps

Q&A



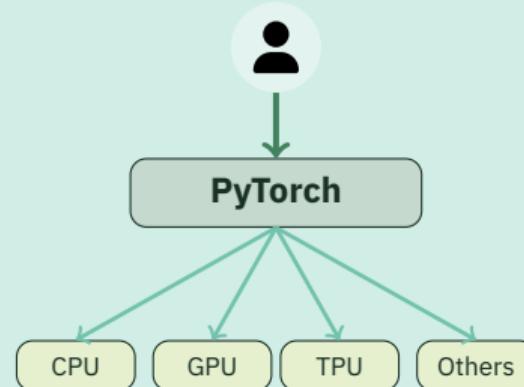
The Infrastructure Challenge



What is PyTorch?

A Leading Open Source ML Framework:

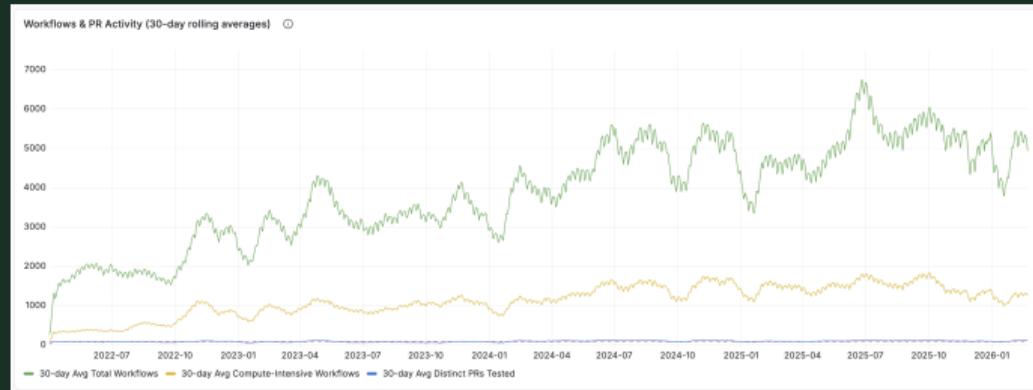
- › Researchers, data scientists, ML engineers
- › Research and production ML workloads
- › (Distributed) Training, LLM, RL, Inference
- › Python API, C++ Core (libtorch)
- › Eager and Graph Mode (Inductor)
- › Rapid growth and adoption across industry



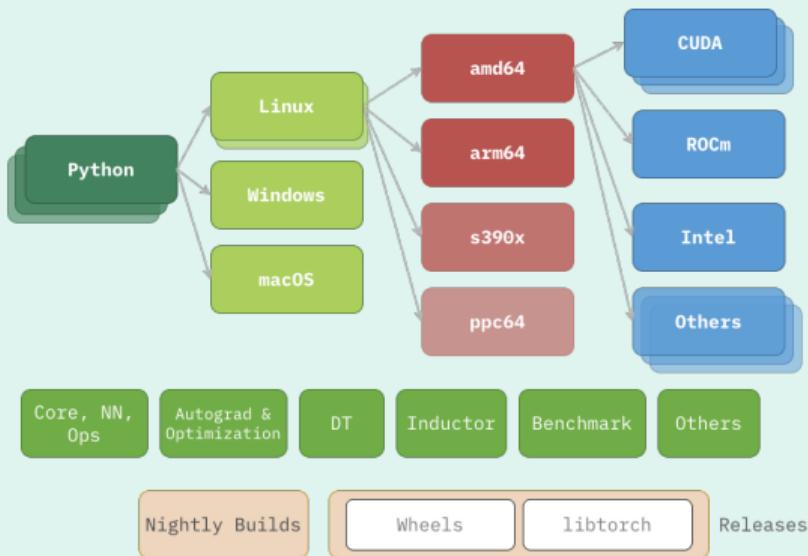
CI/CD Infrastructure Today

- › ~1M\$ monthly infrastructure costs
- › ~1M hours of compute per month
- › ~100 distinct PRs/day
- › ~1.3k Compute-intensive Workflows/day
- › Growing year over year

pytorch/pytorch repo only, numbers are estimated

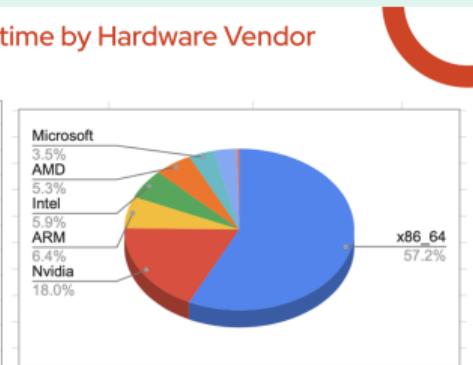


An Expanding Test Matrix



Combined Hardware Runtime by Hardware Vendor
September 2025

Combined Runtime By Hardware Vendor	
Vendor	Total
x86_64	834,085 hours
Nvidia	262,878 hours
ARM	93,000 hours
Intel	86,090 hours
AMD	77,517 hours
Microsoft	51,515 hours
Apple	46,812 hours
IBM	6,245 hours
N/A	0 hours
1,458,142 hours	



* Data from PT Foundation AWS Account + PyTorch HUD. This does not include self-hosted runner costs by member provided runners.



The Challenges

Community Challenges:

- › Maintain high-quality end-user experience
- › Preserve contributor workflow
- › Provide clear path for vendor engagement

Scale Challenges:

- › **Engineering:** Managing a Diverse Fleet, Access to Platforms
- › **Financial:** Reconcile Community Wishes with Budget Constraints

Security Challenges:

- › Trust and Consistency in a diverse environment
- › Central Administration
- › Trusted Builds

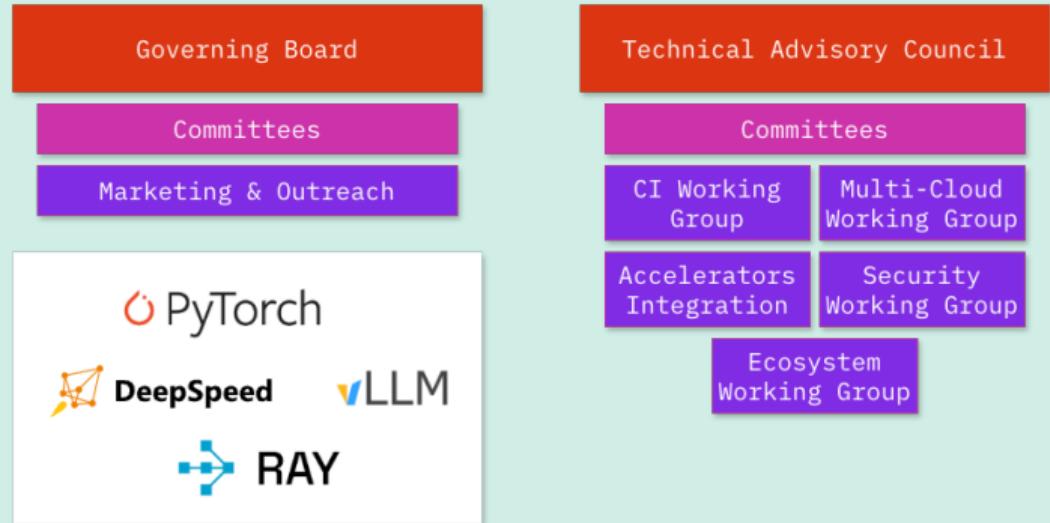


PyTorch Foundation & Working Groups



The PyTorch Foundation

- › Part of the Linux Foundation
- › Neutral home for PyTorch project
- › Hosts multiple projects
- › Thought Leadership (GB)
- › Technical Leadership (TAC)
- › Marketing & Outreach



Working Groups

Multi-Cloud CI

Cloud Agnostic Infrastructure:

- CI/CD jobs and Supporting Infra
- Metrics and Monitoring
- Fleet Management

Vendor-managed Runner Pools

CI Infra

Responsible for the CI/CD infra:

- Develop and maintain tools
- Operate the AWS/GitHub fleet
- Execute Releases

Accelerator Integration

Software Integration:

- Developer Guide
- Framework Improvement

Reference Test Framework

CI Event Relay Infrastructure

Security

- PyTorch Security Triage
- Tools and reports
- CI/CD Security



Governance Model



Photo by Khampha Phimmachak, CC0



Personas: A Balancing Act

End Users

- Stable software
- Available on my dev platform
- Available on my prod platform
- Following the industry at speed

Platform Vendors

- PyTorch for my platform
- Demonstrate platform compatibility
- Low bar to contribution of infrastructure
- Clear guidance and processes

Project Maintainers

- Project Success
- Stay relevant for end-users
- High quality secure builds
- Manageable CI/CD System (scale)
- Smooth CI/CD experience for contributors

PyTorch Foundation

- Ensure Vendor Neutrality
- Vendor participation encouraged
- Fair representation
- Financial sustainability



Integration Models

Three Ways to Contribute Infrastructure:

- › **Public Cloud Credits:** Financial contribution for shared infrastructure
- › **Vendor-Managed Runner Pools:** Tightly integrated with GitHub Actions
- › **Vendor-Managed Test Infrastructure:** Loosely integrated with GitHub

	Community	Scale	Security
Advantages	Flexible models accommodate different vendor capabilities	Outsourcing engineering with vendor-managed setups	More loosely coupled systems cannot compromise core CI system
Challenges	Hide the underlying infra complexity to contributors	Consistent tech stack required for public cloud credits. Consistent metrics and monitoring across the board.	Ensure security best practices on vendor managed infrastructure



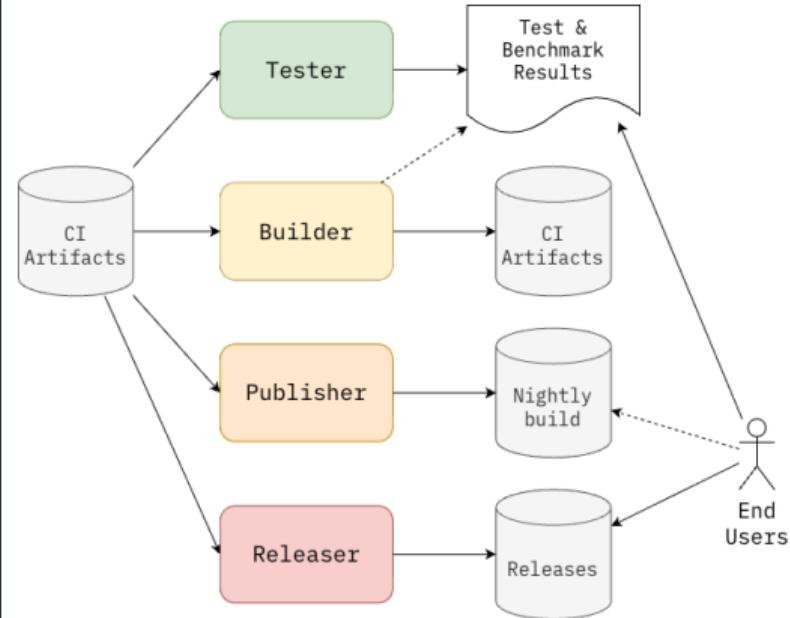
Roles & Requirements

Least privilege access to cloud resources:

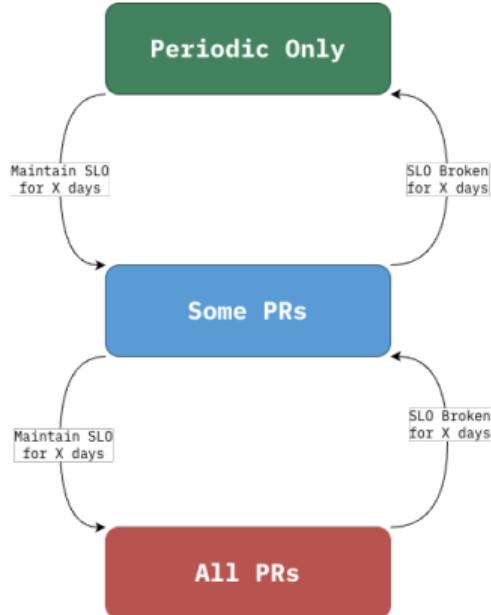
- › Jobs specify nodes access level required
- › Runners are assigned a role
- › Roles define the level of access to resources

Requirements for Runners:

- › Configuration
- › Provisioning
- › Monitoring
- › Security



Triggers & SLOs



Trigger levels for CI/CD jobs:

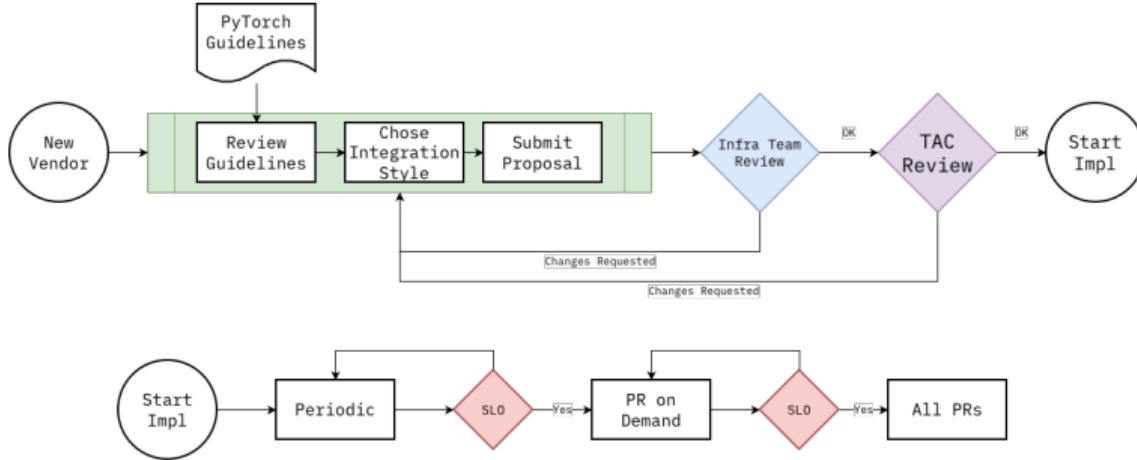
- On-demand only
- Periodic nightly
- Periodic multi-times per day
- Specific PRs only
- All PRs

Service Level Objectives:

- Availability %
- Reliability (Num of infra failures)
- Average Job Queue Time (p95/3 months)
- Engineering Commitment (on-call, slack, meetings)



Contribution Process



- › Working Group provides guidance and recommendations
- › Infra team verify checks for compliance
- › TAC Community for final approval
- › Transparent process with clear requirements
- › Monitoring for SLOs



Central Administration & Visibility

Maintaining Manageability at Scale:

- › Common software stack across clouds
Reusable by vendors too
- › Self-service Tools with Central Admin Override
- › Vendor playbooks, reusable IaaC

Visibility & Control:

- › Public Dashboards: Transparency for community and vendors
- › Cost Tracking: Attribution per vendor and workload type
- › Performance Metrics: Build times, success rates, SLA compliance

A screenshot of a GitHub Actions dashboard for the pytorch/pytorch repository. The timeline shows several recent pull requests (PRs) with their authors, PR numbers, and commit messages. The commits are shown as a grid of colored squares representing different authors.

Time	SHA	Commit	PR	Author
12:39 pm	F77D801	[reverted] Implement RA reader in clang...	#172801	stef-maria
0:54 am	4E7B001	add Python 3.9 support to clang-format	#172802	stef-maria
0:55 am	2D98001	[reverted] Implement RA reader in clang...	#172803	stef-maria
9:13 am	3951001	[revert] Update build requirements comment to #172...	#172804	pranjaygupta
9:13 am	5B85001	[revert] Implement RA reader in clang...	#172805	stef-maria
11:17 am	4E7B001	[reverted] Implement RA reader in clang...	#172806	stef-maria
5:40 am	2D98001	[revert] Implement RA reader in clang...	#172807	pranjaygupta
5:20 am	4E7B001	[revert] Implement RA reader in clang...	#172808	stef-maria
5:20 am	2D98001	[revert] Implement RA reader in clang...	#172809	stef-maria
8:29 am	3212001	[revert] Implement RA reader in clang...	#172810	stef-maria
8:29 am	4E7B001	[revert] Implement RA reader in clang...	#172811	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172812	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172813	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172814	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172815	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172816	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172817	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172818	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172819	stef-maria
8:30 am	4E7B001	[revert] Implement RA reader in clang...	#172820	stef-maria

A screenshot of a GitHub Actions Runner Service dashboard titled "Dashboard". It shows the status of four runners: 0 Active, 4 Offline, and 0 Pending. Below this, there is a section for "Recent Activity" showing three recent events: "batch_update_runner" by admin@githubs.org, "batch_delete_runner" by admin@githubs.org, and "label_policy_update" by alice@githubs.org.

A screenshot of a GitHub Actions Runner Service dashboard titled "Security Events". It lists four security events: "batch_create_runner" (high severity, actor: admin@githubs.org), "batch_delete_runner" (high severity, actor: admin@githubs.org), "label_policy_update" (medium severity, actor: alice@githubs.org, target: alice-runner-d8fc12), and "batch_delete_runner" (high severity, actor: admin@githubs.org).

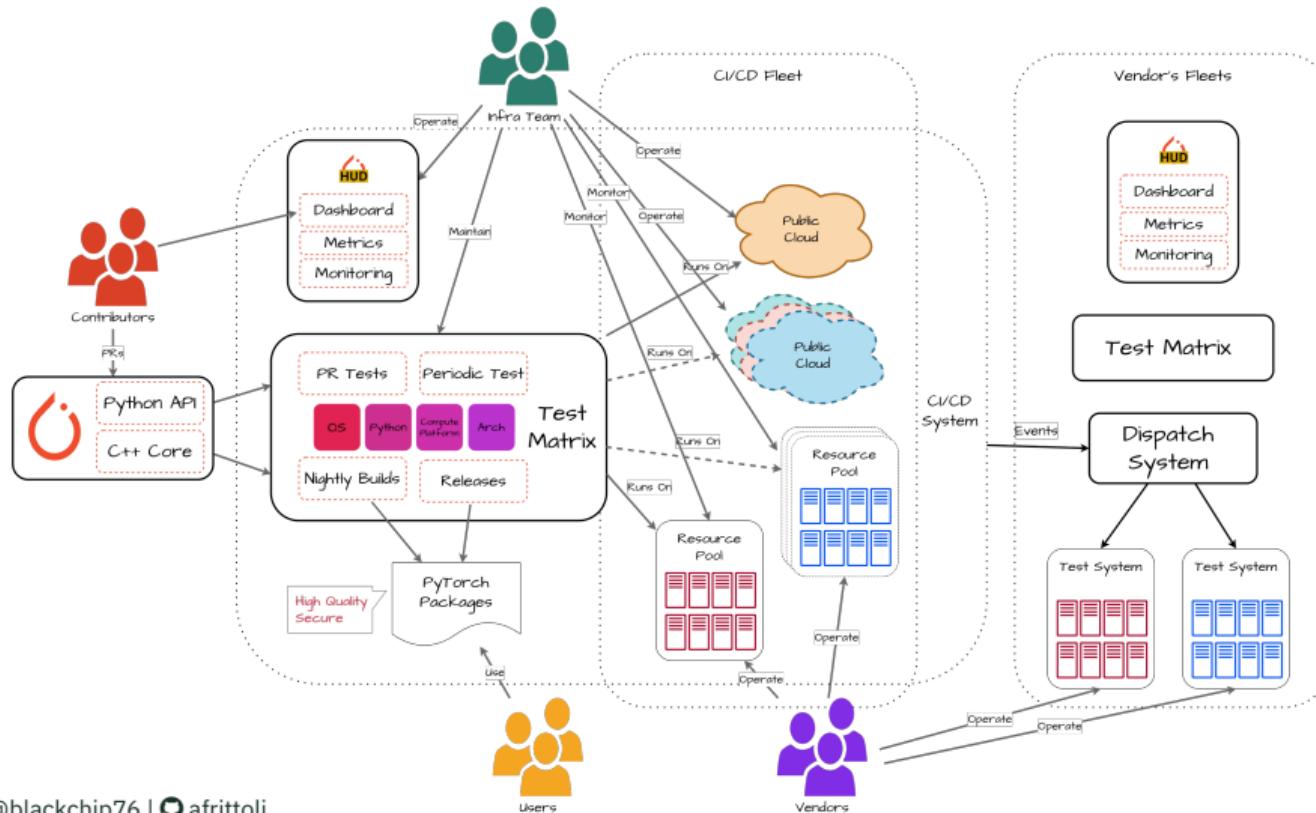
TIMESTAMP	EVENT TYPE	SEVERITY	ACTOR	TARGET
Jan 26, 2020 at 5:09 PM	batch_create_runner	high	admin@githubs.org	-
Jan 26, 2020 at 5:09 PM	batch_delete_runner	high	admin@githubs.org	-
Jan 26, 2020 at 5:40 PM	label_policy_update	medium	alice@githubs.org	alice-runner-d8fc12
Jan 26, 2020 at 6:32 PM	batch_delete_runner	high	admin@githubs.org	-



Technical Architecture



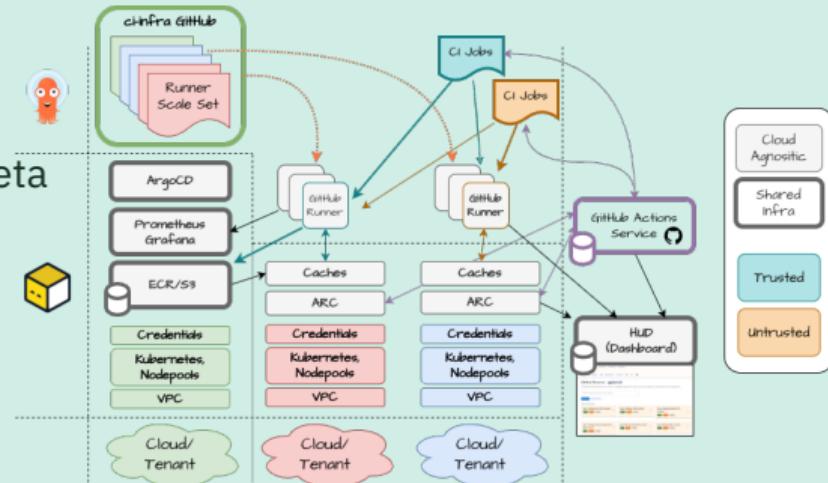
High Level Architecture



Public Cloud Infrastructure

Current AWS Setup:

- > Primary infrastructure hosted on AWS
- > Two accounts, funded by AWS Credits and Meta
- > Self-hosted GitHub Actions runners
- > Managed by PyTorch Infra team
- > Autoscaling based on workload demand



Multi-Cloud & Infra Working Group Initiatives:

- > **Portable CI Jobs:** Run in a container, remove cloud-specific assumptions
- > **Reusable Autoscaler:** Developing portable autoscale based on ARC
- > **Infrastructure as Code:** Reusable OpenTofu modules



Vendor-Managed Runner Pools

Challenges:

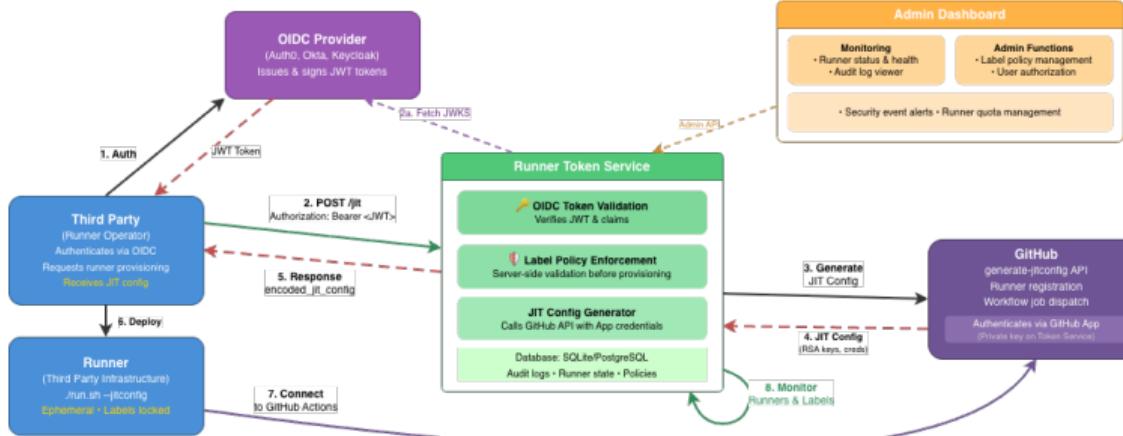
- › Long lived admin credentials
- › Runner metadata cannot be enforced
- › Lack of central administration
- › Lack of auditability

Proposed Solution:

- › GHA Runner Token Service (GHARTS)
- › Use LFID for authentication
- › No Admin Credentials to vendors
- › Ephemeral, non reusable credentials
- › Enforcement of metadata and quota
- › Centralized administration
- › Auditability



GHARTS Provisioning Flow



- Vendor authenticates using existing credentials (LFID)
- Vendor requests a JIT config using the JWT
- GHARTS verifies AuthN and AuthZ for vendor
- GHARTS requests JIT config from GitHub
- Vendor uses JIT config to provision a runner
- JIT lasts 1h, and can only be used Once



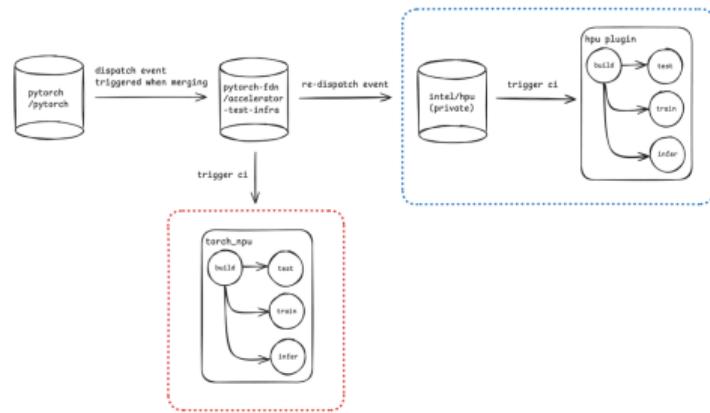
Vendor-Managed Test Systems

Accelerator Integration Working Group Initiative:

- Loosely coupled with PyTorch CI
- Events relayed to external repository
- Vendors provide test infrastructure and environments
- Out-of-tree test workflows

Use Cases:

- Specialized hardware testing (TPUs, custom accelerators)



Old Issues, New Issues

PyTorch Queue Time Analysis

* All datetime values are in UTC. Local: 2026-02-24 21:36:29, UTC Time: 2026-02-25 05:36:29

* Data is collected every 30 minutes, including all jobs in queue at that time. ⓘ

Chart Type

New Challenges:

- > Multiple routes to Integration
- > Address several challenges, but
- > Creates extra complexity
- > Impedes visibility across the board

Ongoing Work to Solve:

- > Collaboration across working groups
- > Same visibility of test results for all routes
- > Consistent metrics
- > Single administrative interface



avg queue time line

Time Range Last 3 Days Granularity Half Hour

Search Category workflow name

By default, shows data for all queued jobs. Using filter and checkbox below for specific items.

Select items:

Filter

Select All Unselect All

- pull
- inductor
- operator_microbenchmark
- inductor-A100-perf-nightly
- inductor-perf-nightly-h100
- windows-binary-libtorch-debug
- trunk
- linux-aarch64-binary-manywheel
- periodic
- inductor-unitest
- windows-arm64-binary-wheel
- trunk-rocm-sandbox
- windows-binary-wheel

Search

* Click to apply filter changes



Results & Lessons Learned



Fleet Composition

Onboarded:

- › **Public Cloud Credits:** AWS, Meta
- › **Vendor-Managed Runner Pools:** AMD ROCm, IBM s390x, IBM ppc64, Intel XPU, NVIDIA CUDA B200, NVIDIA Windows RTX
- › **Vendor-Managed Test Systems:** Huawei Ascend NPUs, Intel Gaudi HPU
- › ...and GitHub Hosted runners too!

In progress:

- › **Vendor-Managed Runner Pools:** IBM s390x (via GHARTS)
- › **Vendor-Managed Test Systems:** RedHat RHEL



Ongoing Work

Public Clouds

- Migrate CI Jobs (with the help of AI!)
- Migrate Autoscaler to ARC based system

Runner Pools

- Deploying GHARTS
- Onboarding IBM pools
- Integrate ARC with GHARTS

Test Systems

- Integrate into hud
- Enable PR comments and votes

Governance

- Finalize onboarding process
- Finalize SLO definitions



Key Lessons Learned

Principles for Collaborative Infrastructure:

- › **Fair Governance:** Protect everyone's interests, vendors can't veto but have voice
- › **Engage Vendors Early:** Let them help build solutions, not just consume them
- › **Empower Maintainers:** Give infra team authority and tools to manage at scale
- › **Security First:** Build security in from day one, not as afterthought
- › **Flexible Models:** Multiple contribution paths reduce friction and attrition

Challenges:

- › Balancing vendor needs with project sustainability
- › Standardization across diverse platforms and requirements



Applicability to Other LF Projects

Shared Challenges:

- › Accepting vendor infrastructure
- › Maintaining project neutrality
- › Ensuring transparency
- › Achieving financial sustainability

This Model Can Help:

- › Proven governance framework
- › Technical architecture patterns
- › Onboarding playbook



Future Directions

Ongoing Activities:

- › Complete GHARTS rollout with IBM and additional vendors
- › Enhance monitoring and observability across all integration models
- › Expand portable CI/CD tooling for multi-cloud deployments
- › Standardize event streaming with OpenTelemetry/CDEvents

Long-term Goals:

- › Share governance and technical patterns with other LF projects
- › Build community of practice around collaborative infrastructure
- › Continuously refine based on vendor and maintainer feedback



Key Takeaways

- › **Sustainable Scaling:** Distribute infrastructure costs while maintaining project control
- › **Governance Matters:** Fair rules enable vendor participation without compromising autonomy
- › **Security & Transparency:** Non-negotiable foundations that build trust
- › **Flexible Models:** Multiple contribution paths reduce friction and increase adoption
- › **Reusable Pattern:** This model can help other Linux Foundation projects





Questions?



References & Contact

Resources:

- › PyTorch Foundation: <https://pytorch.org/foundation>
- › PyTorch CI/CD: <https://github.com/pytorch/pytorch/wiki/CI>
- › GitHub: <https://github.com/pytorch/pytorch>
- › Multi-Cloud Infrastructure WG: PyTorch Slack #infra-wg

Contact:

- › Andrea Frittoli
- › andrea.frittoli@uk.ibm.com
- ›  @blackchip76 |  afrittoli

